The Nixon Corpus: Speaker-Annotated Transcripts of the Nixon Tapes Using a Fine-Tuned Transformer Model

Arvin Akoopie

Preprint, version 1

Abstract

Richard Nixon's secret White House taping system recorded 2,719 hours of conversation between February 16, 1971 and July 18, 1973. Hand transcription reached only the conversations needed for trials, diplomatic histories and edited selections. All 19,426 conversations released as individual recordings were transcribed into the Nixon Corpus: 2,206 hours, 19.0 million words, every word timed and every speech segment labeled. The first publicly documented speech recognizer fine-tuned on the tapes, a Transformer model (NVIDIA's Parakeet TDT 0.6B), was paired with a second, independent recognizer (Whisper large-v3), and a large language model was used to reconcile the two readings against the National Archives' subject logs, H. R. Haldeman's diary and the day's other conversations. Every change was recorded as a reversible edit. Fourteen principal speakers were named by a voiceprint classifier. On 165 held-out conversations transcribed by the State Department's historians, the corpus recovered 76.1% of words and 80.1% of names, against 69.2% and 60.0% for Whisper alone. On the 130 of them outside the classifier's training data, 92.4% of named words were attributed to the right speaker. Transcripts, model and code were released under open licenses.

Keywords: Nixon White House tapes; presidential recordings; archives; speech recognition; speaker identification; large language models; Watergate

1 Introduction

On February 16, 1971, the Secret Service, at Nixon's request, began recording his conversations in the Oval Office and the Cabinet Room. Within sixteen months the system covered his office in the Executive Office Building (EOB), telephones in three rooms and Camp David. The collection ended on July 18, 1973: 2,719 hours on 4,042 reels (Nixon Library, 2019), cataloged by the National Archives and Records Administration (NARA) as 22,723 conversations (NARA, n.d.) and released in stages until August 21, 2013 (Nixon Library, 2019). It was an unusually complete record of private deliberation in a modern presidency: the opening to China, the end of the Vietnam War, the 1972 election and, conversation by conversation, Watergate.

Almost none of it had been transcribed. The Watergate Special Prosecution Force (WSPF) and the House Judiciary Committee transcribed what the prosecutions and impeachment needed (Watergate Special Prosecution Force, 1973–1975; U.S. House of Representatives, Committee on the Judiciary, 1974); the State Department's *Foreign Relations of the United States* (FRUS) printed foreign-policy conversations (U.S. Department of State, n.d.); scholarly editions selected (Kutler, 1997; Brinkley and Nichter, 2014, 2015; Presidential Recordings Digital Edition, n.d.; Dean, 2014). Everything else had only NARA's subject logs, outlines of each conversation's topics (Nixon Library, n.d.). A registry of machine-readable published transcripts, compiled for this work, covered 426 conversations: 1.9% of the catalog.

The digitized tapes have heavy amounts of static and background noise, which make transcription difficult. Microphones hidden in the Oval Office desk, beside its fireplace and in the other rooms, and taps on the telephones recorded at 15/16 inches per second, amid overlapping speakers and room noise (Nixon Library, 2019). None of the earlier AI-assisted transcriptions (Fine, 2026; Hoover Institution, 2026; Ferguson and Nichter, 2026) documented adapting a recognizer to the tapes. The most widely used recognizer, Whisper, was designed to condition each word on its own previous output, and it was documented to loop, skip speech and write text over silence (Radford et al., 2023; Koenecke et al., 2024; Barański et al., 2025; Bain et al., 2023). Parakeet TDT, a FastConformer (a Transformer with convolution) with a transducer decoder, was designed to predict each word chiefly from the audio (Xu et al., 2023; NVIDIA, 2025).

Both recognizers were therefore used. A large language model was given their two independent readings and the archive's own records, and it returned auditable edits; earlier systems voted between recognizers (Fiscus, 1997) or rewrote one recognizer's hypotheses (Chen et al., 2023). A classifier was trained to name the principal speakers instead of clustering anonymous voices (Bredin, 2023). The contributions of this work were the first publicly documented speech recognizer fine-tuned on the tapes, a Transformer model; a two-recognizer pipeline with archival correction; speaker labels for fourteen voices; a held-out evaluation on 165 FRUS transcripts, with a corpus-wide measurement of single-model artifacts; and the open corpus, model and code.

2 Data and methods

Each conversation was transcribed by two recognizers: Parakeet, fine-tuned on the tapes, and Whisper, unchanged. The two outputs were merged: Parakeet's words were kept and Whisper's reading was attached at every disagreement. A large language model was then used to correct the merged text against the archive, returning a list of edits. Speakers were labeled by a voiceprint classifier (Figure 1). Every layer was kept in the release.

2.1 Sources

Audio. 19,426 conversations, February 16, 1971 to July 12, 1973, 2,206 hours, were obtained from NARA's catalog copies and the Nixon Library's older per-conversation files: 6,548 telephone conversations, 6,224 in the Oval Office, 4,435 in the EOB, 2,003 at Camp David and 216 in the Cabinet Room. These made up 81.1% of the 2,719 hours in the Nixon Library's finding aid; 44 further conversations could not be fetched, and 3,253 were available only within whole-tape files. NARA's withdrawal tones (1,242 spans in 841 conversations, 18.4 hours) were muted before transcription.

Archive. The archival sources were NARA's conversation manifest (date, time, location, participants); the subject logs, 2,383 recovered from the Library's truncated archive and matched to 94% of conversations; and H. R. Haldeman's diary, 1,473 entries from January 18, 1969 to April 30, 1973 (Haldeman, 1994; Nixon Library, 2018).

Reference transcripts were divided into two strictly separated sets. *Train:* everything used to build and tune the system: 60 WSPF transcripts (393,443 words) for fine-tuning, and eleven references of eight conversations (six room recordings, two telephone calls) that were used to set the learning rate, the correction instructions and the confidence threshold. These eight were not in the fine-tuning data; five of their room recordings, the *train room recordings*, were among the corpus conversations. *Test:* every FRUS document whose editors

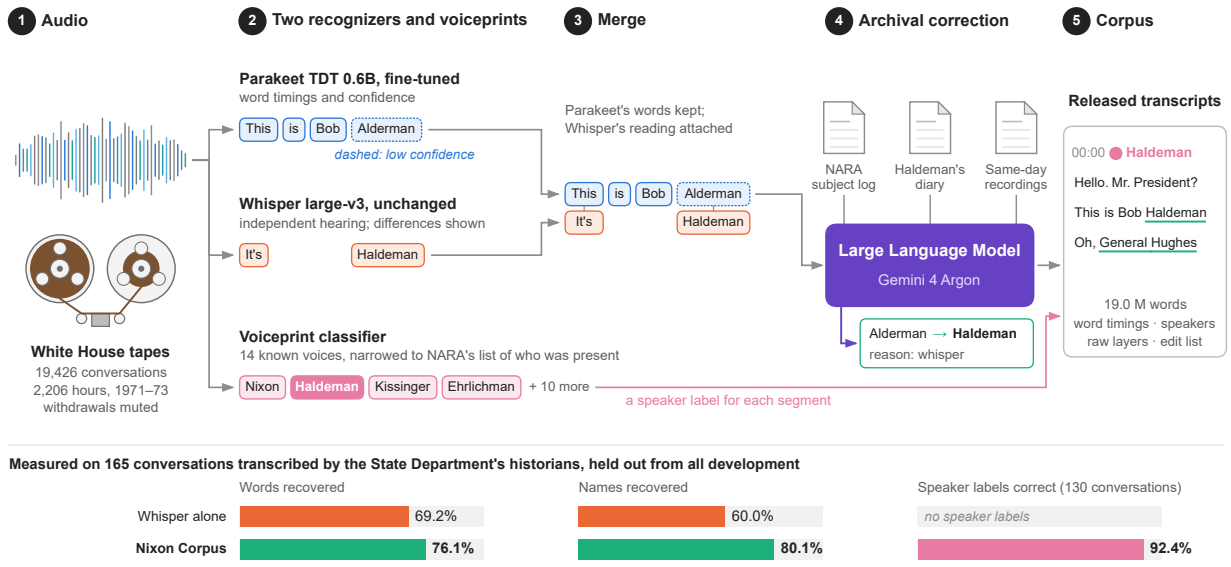


Figure 1: The pipeline, followed through one moment of the corpus: Haldeman giving his name on the telephone (018-078, January 13, 1972). Parakeet was unsure of *"Alderman"* (dashed); Whisper heard *"It's Bob Haldeman"* (only its readings that differed from Parakeet's are shown); the large language model adopted Whisper's reading and cited it; the classifier picked Haldeman among its fourteen voices. Edits are underlined in the released transcript. Bottom: accuracy on the 165 held-out test conversations (speaker labels: the 130 outside the classifier's training data).

transcribed one conversation from the tape as speaker-labeled dialogue, each excerpt aligned on its own; after exclusions, 165 conversations, February 26, 1971 to June 18, 1973, 351,883 reference words (119 Oval Office, 26 telephone, 10 EOB, 5 Cabinet Room, 5 Camp David). No FRUS text reached the correction model. *Registry:* NARA's WSPF and NARA transcript lists, the train references and these FRUS documents, 426 conversations. For the passages of Section 5, both Brinkley and Nichter volumes and Dean (2014) were also searched; Kutler (1997) and the Presidential Recordings Digital Edition were not.

2.2 Two recognizers

Parakeet, fine-tuned. Parakeet TDT 0.6B v3, a 600-million-parameter FastConformer (a Transformer with convolution) with a token-and-duration transducer decoder (Xu et al., 2023; NVIDIA, 2025), was chosen because it was not conditioned on its own output; on the corpus it ran at least 4.3 times faster than Whisper. It was fine-tuned on 4,790 WSPF clips (19.9 hours, 49 conversations) force-aligned with wav2vec 2.0 (Baevski et al., 2020), plus labels the author made by ear in twelve hours: 68 transcribed clips, 98 corrected chunks of low-confidence room audio, 38 noise clips on which the model had to write nothing, and 22 stretches marked with the speaker's name, each repeated ten times. Room recall rose from 63.8% to 68.5% (WSPF clips), 70.0% (plus transcribed clips), 70.7% (plus corrected chunks) and 72.3% for the released model, v5 (plus noise clips; learning rate 6×10^{-5}, AdamW, 8 epochs, NVIDIA NeMo 3.0.0, Kuchaiev et al., 2019). Decoding the corpus cost US$8.60 of rented GPU time; 88.4% of Parakeet's 18,274,746 words received a confidence score.

Whisper and the merge. Whisper large-v3, unchanged, transcribed 19,421 conversations, a first fifth (3,891) in batches with WhisperX (settings in the repository). The two outputs were aligned word by word. Agreed and Parakeet-only words were kept; Whisper's reading was attached at every disagreement; Whisper's words

filled Parakeet's silences (no Parakeet word within 0.5 seconds, mean Whisper probability at least 0.5); Parakeet words below confidence 0.01 were marked.

2.3 Correction against the archive

Recognizers often turned names into ordinary words: Parakeet wrote *"the whole different"* for Haldeman and *"Generally"* for John Ehrlichman where Whisper heard the names. The correct name was often in the archival record. Google's Gemini 4 Argon (Google DeepMind, 2026) (batch interface, temperature 0.2, October 20 to 23, 2026) was used to reconcile the two recognizers' readings with NARA's record and subject log and Haldeman's diary for that day and the next, in order to create a single corrected transcript per conversation. The merged transcript was sent in 60-second blocks, with Whisper's alternatives inline and low-confidence words marked. Short conversations of the same day were sent in one request, so that one recording could correct another. The model was instructed to keep fillers, repetitions and grammar, never to paraphrase, and to mark nonsense as `[unclear]` (instructions: `src/argon.py`). It returned edits rather than text: block, words replaced, replacement, confidence and the main evidence (`whisper`, `nara`, `diary`, `day`, `context`, `knowledge` or `sound`). 98.5% of edits were located and applied by exact match; every edit was marked and kept reversible. All conversations with words were corrected except six. The pass used 57.3 million input and 19.4 million output tokens and cost US$154.13 at the batch price (US$0.07 per audio hour), or US$308.26 at the standard price at which Google billed the batches completed by October 21.

2.4 Speakers

Speakers were named by a logistic-regression classifier on frozen WeSpeaker ResNet-34 embeddings (Wang et al., 2023), normalized by room type, over fourteen voices (Nixon, Kissinger, Haldeman, Ehrlichman, Colson, Dean, Ziegler, Mitchell, Rogers, Shultz, Haig, Woods, Connally, Butterfield) or "other". It was trained on 24,656 clips from 902 conversations: WSPF speaker turns, two-person recordings on which Nixon was confidently ruled out (bringing in Kissinger, Haig, Rogers and others), and the author's 22 marks. Each Whisper segment of at least 1.5 seconds was scored among the people NARA listed as present and was named at probability 0.7 or more.

2.5 Evaluation

Hypotheses and references were normalized with Whisper's English normalizer (Radford et al., 2023), cut to their shared stretch and aligned word by word. The measures were word recall, $1-(S+D)/N$; word error rate, $(S+D+I)/N$; name recall and name precision over the reference's names (words capitalized mid-sentence at least twice and in at least 80% of their uses; a wrong name absent from the reference was not counted); and recall on Nixon's turns and on replies of four words or fewer. Test intervals were 95% bootstrap intervals over conversations, paired between layers. The test set informed no decision and shared no conversation with the train set. Of 231 FRUS conversations, 66 were excluded: 44 not in the corpus (43 released only within whole-tape files, one released after the conversation list was built), 9 the author had heard, labeled or corrected, 7 parallel recordings of excluded conversations and 6 whose audio held a different conversation. Because FRUS text could have been in the correction model's pretraining data, the merged layer, untouched by any language model, was scored beside the corrected layer.

Table 1: Accuracy of each corpus layer (%). Train: the five train room recordings (48,189 reference words). Test: 165 held-out FRUS conversations (351,883 reference words), 95% bootstrap intervals over conversations.

	Train (five room recordings)			Test (165 FRUS conversations)		
Layer	Word recall	WER	Names	Word recall	WER	Names
Whisper large-v3 alone	57.2	45.8	48.8	69.2 [67.4, 70.7]	42.0	60.0 [57.0, 62.8]
Parakeet v5 alone	72.4	31.3	73.0	73.7 [72.0, 75.1]	38.1	60.2 [57.2, 63.2]
Merged (no language model)	72.5	31.1	73.1	73.7 [72.1, 75.2]	37.9	60.3 [57.2, 63.2]
Corrected	**73.9**	**29.6**	**81.2**	**76.1** [74.5, 77.6]	**35.1**	**80.1** [77.3, 82.5]

3 Results

3.1 Accuracy

Fine-tuning worked. Word recall on the six train room recordings, none of them in the fine-tuning data, rose from 63.8% to 72.3% (paired gain 8.5 points, 95% interval 5.0 to 10.8; Figure 2c). Nixon's own words rose from 49.3% to 64.9%, short replies from 32.2% to 46.7%, and WER fell from 39.0% to 31.3%.

The corpus beat Whisper on held-out data. On the 165 FRUS conversations the corrected layer recovered 76.1% of reference words; Whisper alone recovered 69.2% (Table 1, Figure 2). The 6.9-point gain (6.3 to 7.7) split into 4.5 points from fine-tuned Parakeet and 2.4 from correction. The corpus won on 160 of the 165 conversations. The recognizer gap was largest on Nixon's own words and on short replies (Parakeet v5 71.6% and 50.8%, Whisper 63.6% and 36.0%).

Correction supplied the names. Parakeet v5 and Whisper recovered 60.2% and 60.0% of names. Correction lifted name recall by 19.8 points to 80.1%, with no loss of precision: 90.3% of displayed reference names were right, against 89.0% for Whisper alone. By location the corrected layer recovered 85.6% of words at Camp David, 83.7% on the telephone, 76.6% in the Oval Office, 76.5% in the Cabinet Room and 49.3% in the EOB (Figure 2d).

3.2 Artifacts of single-model transcription

Whisper alone looped, dropped speech and wrote over silence (Figure 3). Its output held 2,613 repetition loops in 1,463 conversations (55,140 words); Parakeet v5 had 105, the released corrected layer 265. Run in batches with WhisperX (Bain et al., 2023), the configuration of earlier machine transcriptions of the tapes (Fine, 2026), it discarded speech: wherever WhisperX wrote nothing for 20 seconds or more, Parakeet transcribed 5.2% of its words, in 2,869 stretches at 60 or more words a minute. Across NARA's muted withdrawals, Whisper wrote 1,882 words in 422 conversations, typically "you" every 30 seconds, and across the corpus 987 stock phrases such as "thank you for watching"; Parakeet wrote 131 words over the same silence. No word inside a muted passage was kept in the release. Each recognizer also heard words the other missed, and both were therefore kept in the corpus.

3.3 What the second reading and the correction added

At least one recognizer was right on 79.5% of test reference words, 6.0 points above Parakeet alone on the same words; correction recovered 2.4 of those points. **The second reading drove the gain.** On the train room recordings, correction gained 405 reference words from Parakeet's text alone, 564 with Whisper's readings

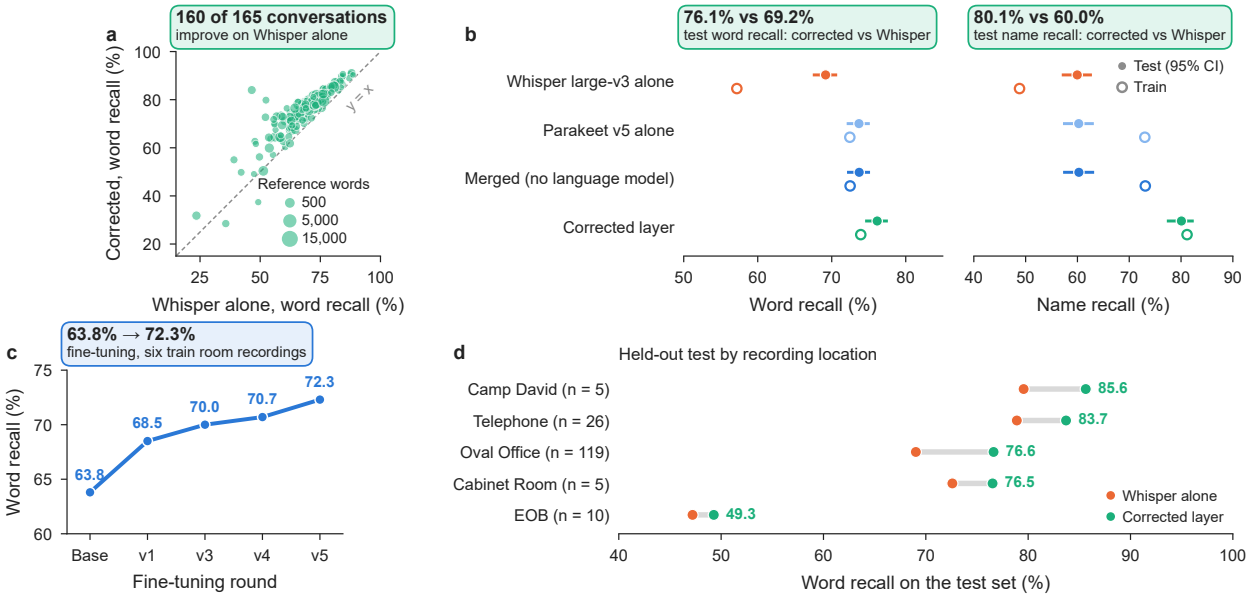


Figure 2: Accuracy. **a**, Word recall of each of the 165 held-out FRUS conversations, Whisper alone against the corrected layer (marker size: reference words). **b**, Word and name recall of each layer on the test set (filled, 95% intervals) and the train room recordings (open). **c**, Fine-tuning: pooled word recall on the six train room recordings, by round. **d**, Test word recall by recording location, Whisper alone and corrected.

and 712 with low-confidence words also marked; before the corpus run, 340 edits helped and 20 hurt. **Same-day context corrected across recordings.** On April 9, 1971, Parakeet's *"want Dr. Burns, uh, Marjorie Burns? I'm pleased to call"* (476-021) became *"Find Dr. Burns, Arthur Burns, I placed a call"*, the words of the telephone recording of the same moment (001-061); the author confirmed both by ear. **Edits cited their evidence**: corrected test-set names cited Whisper's reading (42.8%) or the subject log (34.8%), and the model's own knowledge for only 4.6%.

3.4 Speakers

On the 130 test conversations outside its training data (220,664 aligned words), the classifier named 68.4% of words, and **92.4% of named words carried the speaker named in the FRUS transcript**; for the fourteen voices it was trained on, 95.9% (Nixon 94.9% over 78,305 named words, Kissinger 97.5% over 50,743; Figure 4c). Cross-validated balanced accuracy was 87.2%.

4 The corpus

Release 1.1 (Table 2, Figure 4) was packaged as one JSON record per conversation: NARA's metadata and subject log, audio sources, withdrawn spans, speaker segments, and every word with its timing, speaker, confidence and source (Parakeet, both recognizers, a Whisper fill, a model edit or a correction by ear). Each edit was stored with what it replaced and why. Plain-text transcripts and the raw recognizer outputs were included. A browsing application, included in the repository, was built to present the corpus by day, conversation and person and to play each transcript in sync with the audio; no language model wrote its text.

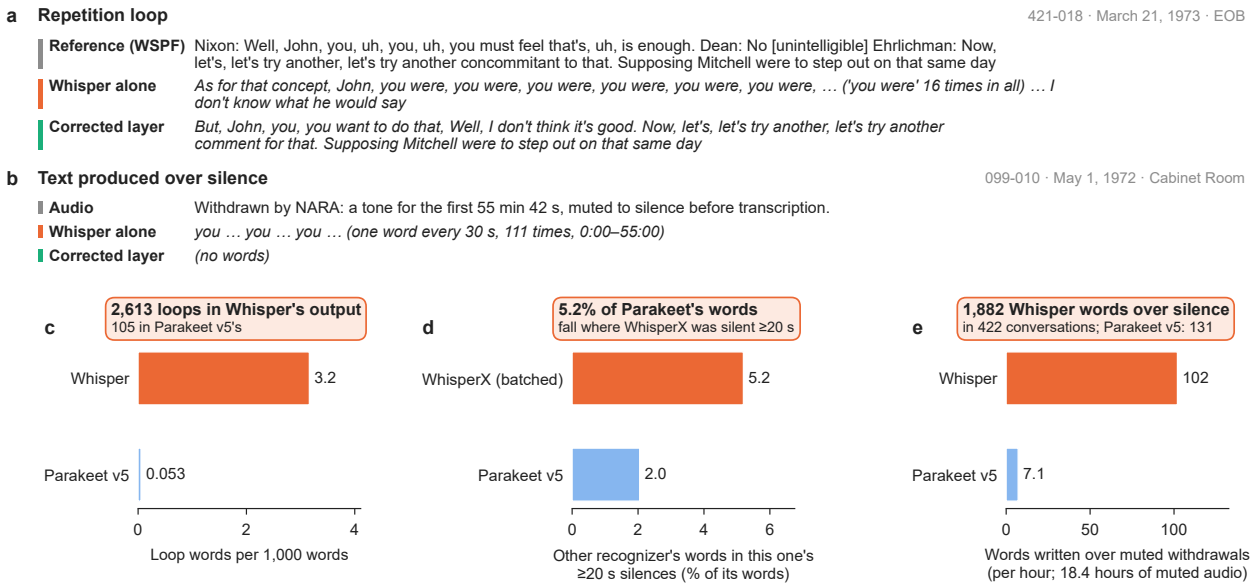


Figure 3: Artifacts of Whisper-only transcription, from the corpus's own Whisper layer. **a**, A repetition loop (421-018): Whisper repeated "you were" sixteen times and lost Ehrlichman's turn; the corrected layer recovered it. **b**, Text over silence (099-010): across a 55-minute withdrawal, muted before transcription, Whisper wrote "you" every 30 seconds; the corrected layer held no words. **c**, Loop words per 1,000 words, whole corpus. **d**, Dropped speech: the other recognizer's words that fell in this one's silences of 20 seconds or more (WhisperX: the batched fifth). **e**, Words written over muted withdrawals per hour of muted audio.

Table 2: The Nixon Corpus, release 1.1.

Conversations	19,426 (February 16, 1971 – July 12, 1973), 17,343 with words
Audio	2,206 hours: 6,548 telephone, 6,224 Oval Office, 4,435 EOB, 2,003 Camp David, 216 Cabinet Room
Words	18,994,348, each with start and end time, source and confidence; every model edit marked and reversible
Speakers	63% of words named; of 1,461 hours of scored speech, 1,023 named (Nixon 637, Haldeman 99, Kissinger 79)
Context	NARA date and location for every conversation, participants for all but 35, subject log for 94%
Format	JSON (gzip), plain text, metadata table, raw layers, checksums; 776 MB

5 Three passages

Every quotation below was taken word for word from the released corrected transcript, with its punctuation and standard capitalization, and was checked by the author against the audio.

5.1 The corpus matched a participant's transcript: 871-004, March 7, 1973

With L. Patrick Gray's confirmation as FBI director going badly (Congressional Research Service, 2018), Nixon told John Dean he wanted friendly senators to question Gray under oath about what J. Edgar Hoover had told him. Dean printed the passage from his own transcription (Dean, 2014, Conversation No. 871-4):

> **Dean (2014):** "I'm going to make him lie, because I think Gray's not handling himself well. What he said yesterday with regard to the fact that this 'jolly well' bullshit … that it was highly improper for you to be present …"

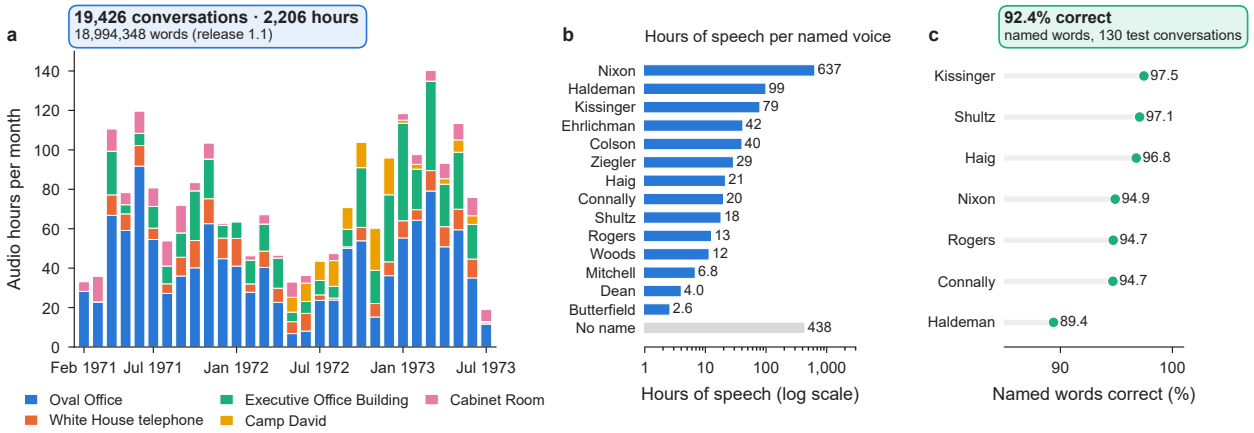


Figure 4: The corpus. **a**, Audio hours per month by recording location, February 1971 to July 1973. **b**, Hours of speech per named voice (segments named at probability 0.7 or more; log scale) and the 438 scored hours without a name. **c**, Named words correct on the 130 held-out test conversations outside the classifier's training data, for voices with at least 200 named words.

> **This corpus (5:26–5:48):** *"I told him that Hoover had told me, and I'm going to make him lie because I think Gray's not handling himself well. What he said yesterday with regard to the fact that, it was, that, you know, this jolly well bullshit and all that sort of thing that you were, that it was highly improper for you to be present."*

These were the recognizers' own words: the correction step made no edit here. Over the surrounding minutes the merged layer shared 95 of its 1,288 five-word sequences with Dean's text, against 10 to 20 for unrelated stretches. Dean printed what he selected; the corpus captured the whole conversation, including Nixon's question at 4:55 about the FBI's 1968 *"electronic surveillance of Vice President Agnew is phone when he was in Phoenix"*. The documented check concerned calls from Albuquerque (Germany et al., 2014; Hughes, 2014): the recognizer heard Nixon correctly, and the slip was his.

5.2 The cover story in Nixon's words: 916-019, May 11, 1973

Eleven days after Haldeman and John Ehrlichman resigned, and days after Vernon Walters testified that they had sought the CIA's help in stopping an FBI investigation (Brinkley and Nichter, 2015), Nixon met Haldeman and Alexander Haig in the Oval Office. At 10:20 Nixon (confirmed by the author's listening) told Haldeman how Ehrlichman's account of the June 23, 1972 meeting with the CIA's leaders should go:

> *"Now, my point is, John's story should be perfect, should be, and I hope you talked to John. John's story is perfect. John should say his only concern was the whole CIA involvement and so forth and wanted to be sure that the, this investigation of Watergate didn't get out, as he put it, didn't get out of hand. […] Getting out of hand had to do with exactly what's now happening. Yeah. Which is unloading all this Pentagon papers. Investigation type stuff, the leak, the plumbers operation type stuff, the, Exactly. Wiretapping of the reporters and, and staff members type stuff."*

"John" was Ehrlichman: the conversation placed *"Ehrlichman and me"* at the meeting and said *"Dean was not there"* and that *"Ehrlichman did do all the talking"*. Parakeet transcribed these words where Whisper wrote nothing; the correction step did not touch them. At 13:00 Nixon named his concern and Haldeman answered (*"paint a"* was corrected by ear from *"faint"* by the author):

Table 3: Further leads, with prior-publication checks. Three other candidates were already in print (Dean, 2001, 2014).

Conversation	Date, place	Passage	Prior publication
789-006	Sept. 30, 1972, Oval Office	*"super apathy among the blacks"*, called *"just exactly where you want it"* before the election	Not in the registry or Brinkley and Nichter (2014)
233-010	Dec. 6, 1972, Camp David	*"make Sweden the black post. I want a black for Sweden"*	Not in Brinkley and Nichter (2014); in earlier machine transcripts (Fine, 2026)
690-011	Mar. 21, 1972, Oval Office	An ambassadorship and a request to *"get at least 150 out of it"*	Not in Brinkley and Nichter (2014); in earlier machine transcripts (Fine, 2026)

> **Nixon:** *"I, I don't need to paint a picture. I just want to be sure that John doesn't indicate that the President directed him to get Walters to get Gray to turn this off and set up a phony story"*
>
> **Haldeman:** *"John would never indicate that. […] What the President tells us to do is not, not public, public consumption."*

It was Nixon who, on June 23, 1972, told Haldeman that the CIA's leaders "should call the FBI in and say that we wish for the country, don't go any further into this case, period!" (Watergate Special Prosecution Force, 1974). No published transcript of this exchange[1] was found in the registry, in either Brinkley and Nichter volume, in Dean (2014) or on the Nixon Library's and Miller Center's pages for the conversation; Kutler (1997) and the Presidential Recordings Digital Edition were not checked.

5.3 The ordinary record, and further leads

Most of the tapes were routine, and the corpus captured that too. On March 14, 1972, after John Mitchell's ITT testimony, Nixon telephoned Charles Colson and tried out a line, *"this is a smear a day Congress rather than a law day of Congress"*; Colson laughed, and they agreed to use it (021-068). No subject log recorded such exchanges. Further leads were surfaced by ranking corpus windows outside the registry for politically charged terms; each was heard and confirmed by the author (Table 3).

6 Limitations and conclusion

Limitations. The corpus was built as a machine transcript and a finding aid: it recovered 76.1% of test reference words, 49.3% in the EOB, and every quotation should be checked against the audio, to which every word was linked. FRUS editors chose the test excerpts; a test set drawn at random from the corpus would measure the whole record. Correction depended on a proprietary model; its edits were marked and kept reversible. Voices without a class were sometimes named (16.9% of their test words). Withdrawn passages were kept muted, and passages about private persons call for care in republication.

Conclusion. Two recognizers that erred differently, reconciled against the archive's own records, turned the released Nixon tapes into a searchable, speaker-labeled text of 19.0 million words. On held-out conversations it recovered 76.1% of words and 80.1% of names, against 69.2% and 60.0% for Whisper alone. It was built to tell a historian where to listen, with measured accuracy, word timings and speaker labels that neither the subject logs nor earlier machine transcripts provided.

[1] Easy Nixon's unreviewed machine transcript of 916-019 (Fine, 2026) did not contain it.

Declarations

Data, code and model availability. The corpus (release 1.1) is archived at `https://doi.org/10.5281/zenodo.23657174`, with a data card that documents its format, coverage and limitations. The code for the pipeline, the evaluation and the browsing tool, the scripts that rebuild the public-domain reference transcripts, and the correction instructions (`src/argon.py`) are at `https://github.com/akoopie/nixon-corpus`. The fine-tuned recognizer is at `https://huggingface.co/akoopie/nixon-parakeet-tdt-0.6b-v5`.

License and citation. The transcripts, annotations and model weights are released under CC BY 4.0; attribution should take the form of a citation of this paper and of the release. The model derives from NVIDIA's Parakeet-TDT-0.6B-v3 (CC BY 4.0; NVIDIA, 2025). The code is MIT-licensed. The recordings, NARA's metadata and the government transcripts are public domain; please credit NARA and the Richard Nixon Presidential Library and Museum as their source.

Use of AI tools. AI systems were used as components of the pipeline (the speech recognizers, the speaker-embedding network and the correction model, Gemini 4 Argon, as reported above) and as assistants: the software was written, the analyses were run and the manuscript was drafted with the help of an AI coding assistant (Claude, Anthropic). The author designed the study, directed all analyses, made the research decisions, performed all annotation and listening, and takes full responsibility for the content of this paper.

Acknowledgments. The author thanks the National Archives and Records Administration and the Richard Nixon Presidential Library and Museum, whose archivists preserved, reviewed, described and released the recordings, and whose subject logs and conversation manifest made this work possible.

Competing interests and funding. The author declares no competing interests. The work received no external funding; computing costs were paid by the author.

References

Alexei Baevski, Henry Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Advances in Neural Information Processing Systems 33*, 2020. URL `https://arxiv.org/abs/2006.11477`.

Max Bain, Jaesung Huh, Tengda Han, and Andrew Zisserman. WhisperX: Time-accurate speech transcription of long-form audio. In *Proceedings of Interspeech 2023*, pages 4489–4493, 2023. doi: 10.21437/Interspeech.2023-78.

Mateusz Barański, Jan Jasiński, Julitta Bartolewska, Stanisław Kacprzak, Marcin Witkowski, and Konrad Kowalczyk. Investigation of Whisper ASR hallucinations induced by non-speech audio. In *ICASSP 2025 – IEEE International Conference on Acoustics, Speech and Signal Processing*, 2025. URL `https://arxiv.org/abs/2501.11378`.

Hervé Bredin. pyannote.audio 2.1 speaker diarization pipeline: Principle, benchmark, and recipe. In *Proceedings of Interspeech 2023*, pages 1983–1987, 2023. doi: 10.21437/Interspeech.2023-105.

Douglas Brinkley and Luke A. Nichter. *The Nixon Tapes: 1971–1972*. Houghton Mifflin Harcourt, Boston, 2014.

Douglas Brinkley and Luke A. Nichter. *The Nixon Tapes: 1973*. Houghton Mifflin Harcourt, Boston, 2015.

Chen Chen, Yuchen Hu, Chao-Han Huck Yang, Sabato Marco Siniscalchi, Pin-Yu Chen, and Eng Siong Chng. HyPoradise: An open baseline for generative speech recognition with large language models. In *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023. URL https://arxiv.org/abs/2309.15701.

Congressional Research Service. FBI director nominations, 1973–2017. Report R44842. https://www.congress.gov/crs-product/R44842, 2018.

John W. Dean. *The Rehnquist Choice: The Untold Story of the Nixon Appointment That Redefined the Supreme Court*. Free Press, New York, 2001.

John W. Dean. *The Nixon Defense: What He Knew and When He Knew It*. Viking, New York, 2014.

Niall Ferguson and Luke Nichter. The Nixon tapes, unlocked. *The Free Press*, September 30. https://www.thefp.com/p/the-nixon-tapes-unlocked-niall-ferguson-luke-nichter, 2026.

Todd Fine. Easy Nixon: Accessible archive of the Richard M. Nixon white house tapes. https://www.easynixon.org/, 2026. Accessed October 21, 2026.

Jonathan G. Fiscus. A post-processing system to yield reduced word error rates: Recognizer output voting error reduction (ROVER). In *1997 IEEE Workshop on Automatic Speech Recognition and Understanding Proceedings*, pages 347–354, 1997. doi: 10.1109/ASRU.1997.659110.

Kent B. Germany, Nicole Hemmer, and Ken Hughes. Lyndon B. Johnson and Cartha D. "Deke" DeLoach on 13 november 1968, conversation WH6811-04-13733. In *Presidential Recordings Digital Edition: Johnson Telephone Tapes, 1968*. University of Virginia Press, Charlottesville, 2014. URL https://prde.upress.virginia.edu/conversations/4006140.

Google DeepMind. Gemini 4 argon: Model card. https://deepmind.google/models/model-cards/gemini-4-argon/, 2026. Published October 13, 2026.

H. R. Haldeman. *The Haldeman Diaries: Inside the Nixon White House*. G. P. Putnam's Sons, New York, 1994.

Hoover Institution. Hoover applied history working group announces landmark transcription of Nixon white house tapes. Press release, September 30. https://www.hoover.org/news/hoover-applied-history-working-group-announces-landmark-transcription-nixon-white-house-tapes, 2026.

Ken Hughes. *Chasing Shadows: The Nixon Tapes, the Chennault Affair, and the Origins of Watergate*. University of Virginia Press, Charlottesville, 2014.

Allison Koenecke, Anna Seo Gyeong Choi, Katelyn X. Mei, Hilke Schellmann, and Mona Sloane. Careless Whisper: Speech-to-text hallucination harms. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1672–1681, 2024. doi: 10.1145/3630106.3658996.

Oleksii Kuchaiev, Jason Li, Huyen Nguyen, Oleksii Hrinchuk, Ryan Leary, Boris Ginsburg, Samuel Kriman, Stanislav Beliaev, Vitaly Lavrukhin, Jack Cook, Patrice Castonguay, Mariya Popova, Jocelyn Huang, and Jonathan M. Cohen. NeMo: A toolkit for building AI applications using neural modules. arXiv:1909.09577, 2019. URL https://arxiv.org/abs/1909.09577.

Stanley I. Kutler, editor. *Abuse of Power: The New Nixon Tapes*. Free Press, New York, 1997.

NARA. White house tapes of the Nixon administration, 1971 to 1973. Conversation-level dataset, Data.gov. https://catalog.data.gov/dataset/white-house-tapes-of-the-nixon-administration-1971-to-1973, n.d. CC0; 22,723 conversations.

Nixon Library. H. R. Haldeman diaries collection finding aid (RN-HRHD; NARA identifier 7787364). https://www.nixonlibrary.gov/sites/default/files/2019-03/HRHaldemanDiariesCollectionFindingAid03212018.pdf, 2018.

Nixon Library. White house tapes: Sound recordings of meetings and telephone conversations of the Nixon administration, 2/16/1971–7/18/1973 (RN-WHT; NARA ID 597542), complete finding aid. https://nixonlibrary.gov/sites/default/files/2019-08/White%20House%20Tapes%20Complete%20Finding%20Aid.pdf, 2019.

Nixon Library. White house tapes: Tape subject logs. https://www.nixonlibrary.gov/white-house-tapes, n.d.

NVIDIA. parakeet-tdt-0.6b-v3. Model card, Hugging Face. https://huggingface.co/nvidia/parakeet-tdt-0.6b-v3, 2025. Released August 14, 2025; CC BY 4.0.

Presidential Recordings Digital Edition. The presidential recordings of Richard M. Nixon. Presidential Recordings Digital Edition. Charlottesville: University of Virginia Press. https://prde.upress.virginia.edu/content/nixon, n.d.

Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 28492–28518, 2023. URL https://proceedings.mlr.press/v202/radford23a.html.

U.S. Department of State. Foreign relations of the United States, 1969–1976. Washington, DC: Government Publishing Office. Digital edition, https://history.state.gov/historicaldocuments; TEI source, https://github.com/HistoryAtState/frus, n.d. Office of the Historian.

U.S. House of Representatives, Committee on the Judiciary. *Transcripts of Eight Recorded Presidential Conversations*. U.S. Government Printing Office, Washington, DC, 1974. 93rd Congress, 2nd session.

Hongji Wang, Chengdong Liang, Shuai Wang, Zhengyang Chen, Binbin Zhang, Xu Xiang, Yanlei Deng, and Yanmin Qian. Wespeaker: A research and production oriented speaker embedding learning toolkit. In *ICASSP 2023 – IEEE International Conference on Acoustics, Speech and Signal Processing*, pages 1–5, 2023. doi: 10.1109/ICASSP49357.2023.10096626.

Watergate Special Prosecution Force. Transcripts of Nixon white house tape recordings. Records of the Watergate Special Prosecution Force, Record Group 460, National Archives and Records Administration; per-conversation copies at https://www.nixonlibrary.gov/white-house-tapes, 1973–1975.

Watergate Special Prosecution Force. Transcript of the conversation of June 23, 1972 between President Nixon and H. R. Haldeman (conversation 741-002), trial exhibit 1, *United States v. Mitchell et al.* https://www.nixonlibrary.gov/sites/default/files/forresearchers/find/tapes/watergate/trial/exhibit_01.pdf, 1974.

Hainan Xu, Fei Jia, Somshubra Majumdar, He Huang, Shinji Watanabe, and Boris Ginsburg. Efficient sequence transduction by jointly predicting tokens and durations. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, 2023. URL https://proceedings.mlr.press/v202/xu23g.html.