The Nixon Corpus: Speaker-Annotated Transcripts of the Nixon Tapes Using a Fine-Tuned Transformer Model

Arvin Akoopie

Preprint, version 1

Abstract

Between February 16, 1971 and July 18, 1973, the Nixon White House secretly recorded 2,719 hours of conversations in the Oval Office, the Executive Office Building, the Cabinet Room and Camp David, and on presidential phone calls. Hand transcription reached only the conversations needed for the Watergate trials, diplomatic histories and edited selections. The National Archives (NARA) has digitized and released all the recordings, but an annotated, accurate transcript of all the audio files has not yet been achieved. Here we present the Nixon Corpus: all 19,426 individual recordings transcribed, 2,206 hours, 19.0 million words, every word timed and every speech segment labeled. To achieve this, a fine-tuned speech recognizer model (Parakeet) was combined in tandem with an independent model (Whisper) to generate two transcripts per conversation, while another voiceprint classifier trained on the tapes assigned speakers to each word. A large language model (Gemini 4 Argon) harmonized these data by taking into context the National Archives' subject logs from each respective tape and Chief of Staff H. R. Haldeman's diaries from that day and the following day. On 165 held-out conversations transcribed by the State Department's historians, the corpus recovered 76.1% of words and 80.1% of names, against 69.2% and 60.0% for Whisper alone. The Nixon Corpus can be used to trace how certain policies developed in the Nixon administration. We have identified several potentially novel histories, such as asking for campaign donations for an ambassadorship and conversations surrounding acting FBI Director L. Patrick Gray.

Keywords: Nixon White House tapes; presidential recordings; archives; speech recognition; speaker identification; large language models; Watergate

1 Introduction

On February 16, 1971, the Secret Service, at Nixon's request, began recording his conversations in the Oval Office and the Cabinet Room. Within sixteen months the system had been extended to his office in the Executive Office Building (EOB), to telephones in three rooms and to Camp David. When the system was shut down on July 18, 1973, the collection held 2,719 hours on 4,042 reels (Nixon Library, 2019), which the National Archives and Records Administration (NARA) cataloged as 22,723 conversations (NARA, n.d.) and released in stages until August 21, 2013 (Nixon Library, 2019). It is an unusually complete record of private deliberation in a modern presidency, covering the opening to China, the end of the Vietnam War, the 1972 election and, conversation by conversation, Watergate.

Almost none of it has been transcribed by hand. The Watergate Special Prosecution Force (WSPF) and the House Judiciary Committee transcribed the conversations that the prosecutions and the impeachment inquiry needed (Watergate Special Prosecution Force, 1973–1975; U.S. House of Representatives, Committee on the Judiciary, 1974). The State Department's historians print foreign-policy conversations in *Foreign*

Relations of the United States (FRUS) (U.S. Department of State, n.d.), and scholarly editions publish selections (Kutler, 1997; Brinkley and Nichter, 2014, 2015; Presidential Recordings Digital Edition, n.d.; Dean, 2014). Everything else has only NARA's subject logs, which outline each conversation's topics (Nixon Library, n.d.). A registry of published transcripts in machine-readable form, compiled for this work, covers 426 conversations, or 1.9% of the catalog.

The digitized tapes have heavy amounts of static and background noise, which make transcription difficult. To keep maintenance low, the recorders, Sony 800B machines on extremely thin tape, ran at their slowest speed, 15/16 inch per second, so that one reel held up to six and a half hours; neither the machines nor the microphones were made for recording conversations. The microphones were hidden in the Oval Office desk, beside its fireplace and in the other rooms, and the telephones were tapped, so voices overlapped and competed with room noise (Nixon Library, 2019). None of the earlier AI-assisted transcriptions (Fine, 2026; Hoover Institution, 2026; Ferguson and Nichter, 2026) documents adapting a speech recognizer to the tapes. The most widely used recognizer, Whisper, bases each word partly on the words it has already written, and it is known to fall into repetition loops, skip speech and write text over silence (Radford et al., 2023; Koenecke et al., 2024; Barański et al., 2025; Bain et al., 2023). Parakeet TDT, a FastConformer (a Transformer with convolution) with a transducer decoder, which emits words as it steps through the audio, instead predicts each word chiefly from the audio (Xu et al., 2023; NVIDIA, 2025).

Here we use both recognizers. A large language model receives their two independent readings together with the archive's own records and returns edits that can each be audited; earlier systems vote between recognizers (Fiscus, 1997) or rewrite one recognizer's candidate transcripts (Chen et al., 2023). Instead of sorting voices into anonymous clusters (Bredin, 2023), a classifier names the principal speakers. The contributions of this work are the first publicly documented speech recognizer fine-tuned on the tapes, a Transformer model; a pipeline that corrects two recognizers against the archive; speaker labels for fourteen voices; an evaluation on 165 FRUS transcripts held out from all development, with a corpus-wide measurement of single-recognizer artifacts such as repetition loops; and the open corpus, model and code.

2 Data and methods

Each conversation was transcribed by two speech recognizers: Parakeet, fine-tuned on the tapes (further trained on examples from them), and Whisper, used unchanged. The two transcripts were merged by keeping Parakeet's words and attaching Whisper's reading wherever the two disagreed. A large language model was then used to correct the merged text against the archive, returning a list of edits. Speakers were labeled by a voiceprint classifier (Figure 1). Every layer, meaning each stage's version of the transcript, was kept in the release.

2.1 Sources

Audio. Recordings of 19,426 conversations, 2,206 hours from February 16, 1971 to July 12, 1973, were obtained from NARA's catalog copies and the Nixon Library's older per-conversation files. They comprised 6,548 telephone conversations, 6,224 in the Oval Office, 4,435 in the EOB, 2,003 at Camp David and 216 in the Cabinet Room, and made up 81.1% of the 2,719 hours in the Nixon Library's finding aid. A further 44 conversations could not be downloaded, and 3,261 were available only within whole-tape files; eight of the 19,426 recordings were Library files that NARA's catalog did not list. NARA's withdrawal tones, which mark passages withheld from release (1,242 spans in 841 conversations, 18.4 hours), were muted before

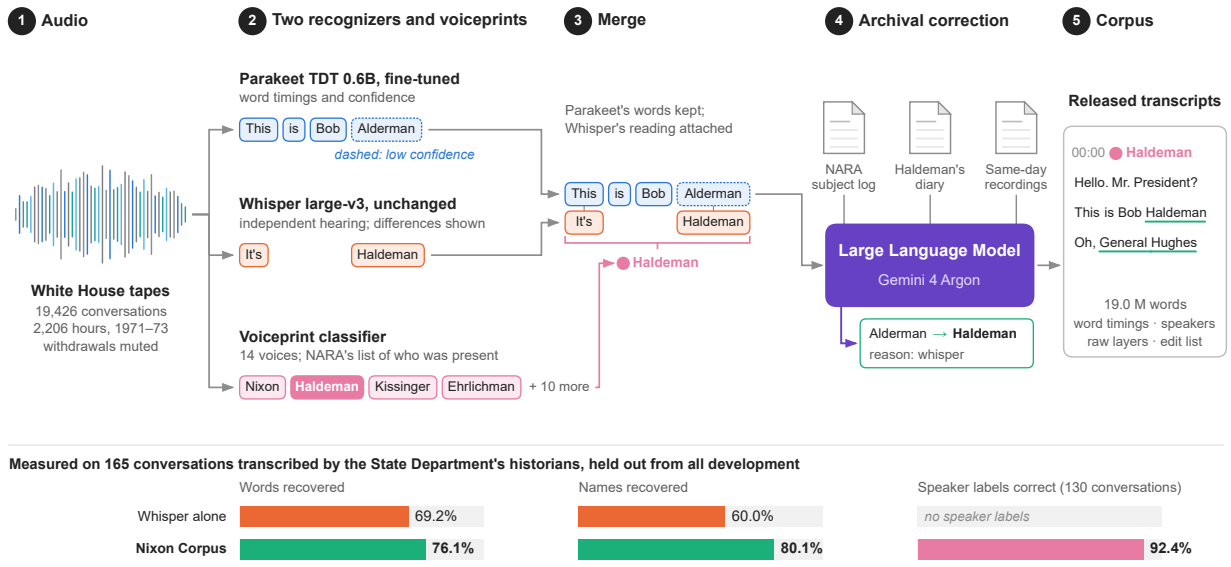


Figure 1: The pipeline, followed through one moment of the corpus: Haldeman giving his name on the telephone (018-078, January 13, 1972). Parakeet was unsure of *"Alderman"* (dashed). Whisper heard *"It's Bob Haldeman"*; only its readings that differed from Parakeet's are shown. The large language model adopted Whisper's reading and cited it, and the classifier picked Haldeman from its fourteen voices and labeled the merged words (bracket). Edits are underlined in the released transcript. Bottom: accuracy on the 165 held-out Test conversations; speaker labels were scored on the 130 of them outside the classifier's training data.

transcription.

Archive. NARA's conversation manifest gave the date, time, location and participants of each conversation. The subject logs, 2,383 of which were recovered from the Library's truncated archive, were matched to 94% of conversations. The diary of Chief of Staff H. R. Haldeman supplied 1,473 entries from January 18, 1969 to April 30, 1973 (Haldeman, 1994; Nixon Library, 2018).

Reference transcripts. These human transcripts, against which the machine output was measured, were divided into two strictly separated sets, Train and Test. The *Train* set held everything used to build and tune the system: 60 WSPF transcripts (393,443 words) used for fine-tuning, and eleven reference transcripts of eight conversations (six room recordings and two telephone calls) used to set the learning rate, the correction instructions and the confidence threshold. These eight conversations were not in the fine-tuning data. Their six room recordings made up the *train room recordings*; five of them were corpus conversations and were used to score the corpus layers. The *Test* set comprised every FRUS document whose editors had transcribed a single conversation from the tape as speaker-labeled dialogue, each excerpt aligned separately. After exclusions it held 165 conversations from February 26, 1971 to June 18, 1973, with 351,883 reference words (119 Oval Office, 26 telephone, 10 EOB, 5 Cabinet Room and 5 Camp David). No FRUS text reached the correction model. The *registry* of published transcripts combined NARA's lists of WSPF and NARA transcripts, the Train references and these FRUS documents, and covered 426 conversations. For the passages of Section 5, both volumes by Brinkley and Nichter and Dean (2014) were also searched; Kutler (1997) and the Presidential Recordings Digital Edition were not.

2.2 Two recognizers

Parakeet, fine-tuned. Parakeet TDT 0.6B v3, a 600-million-parameter FastConformer with a token-and-duration transducer decoder (Xu et al., 2023; NVIDIA, 2025), was chosen because it was not conditioned on its own output, and on the corpus it ran at least 4.3 times faster than Whisper. To adapt it to the tapes, it was fine-tuned on 4,790 WSPF clips (19.9 hours from 49 conversations), whose transcripts were aligned to the audio with wav2vec 2.0 (Baevski et al., 2020), plus labels from the author's twelve hours of listening, each repeated ten times: 68 transcribed clips, 98 corrected chunks of low-confidence room audio and 38 noise clips on which the model had to learn to write nothing. Each addition raised word recall, the share of reference words transcribed correctly, on the room recordings of the Train set: from 63.8% to 68.5% with the WSPF clips, 70.0% with the transcribed clips, 70.7% with the corrected chunks and 72.3% with the noise clips. This last model, v5, was released (learning rate 6×10^{-5}, AdamW, 8 epochs, NVIDIA NeMo 3.0.0, Kuchaiev et al., 2019). Transcribing the corpus cost US$8.60 of rented GPU time, and 88.4% of Parakeet's 18,274,746 words received a confidence score.

Whisper and the merge. Whisper large-v3, used unchanged, transcribed 19,421 conversations; the first fifth of them (3,891) were run in batches with WhisperX (settings in the repository). The two transcripts were then aligned word by word. Words that both recognizers wrote, or only Parakeet wrote, were kept, and Whisper's reading was attached at every disagreement. Where Parakeet was silent (no Parakeet word within 0.5 seconds), Whisper's words filled the gap if their mean probability was at least 0.5. Parakeet words with confidence below 0.01 were marked.

2.3 Correction against the archive

Recognizers often turned names into ordinary words. Parakeet, for example, wrote *"the whole different"* for Haldeman and *"Generally"* for John Ehrlichman, Nixon's domestic policy adviser, where Whisper heard the names. The correct name was often in the archival record. Google's Gemini 4 Argon (Google DeepMind, 2026), a large language model, was used to reconcile the two recognizers' readings with NARA's record and subject log and with Haldeman's diary for that day and the next, producing a single corrected transcript for each conversation. It was run through the batch interface at temperature 0.2 from October 20 to 23, 2026. The merged transcript was sent in 60-second blocks, with Whisper's alternative readings inline and low-confidence words marked. Short conversations from the same day were sent together in one request, so that one recording could correct another. The model was instructed to keep fillers, repetitions and grammar, never to paraphrase, and to mark nonsense as `[unclear]` (see `src/argon.py`). Instead of rewritten text, it returned a list of edits, each giving the block, the words replaced, the replacement, a confidence and the main evidence: `whisper`, `nara`, `diary`, `day`, `context`, `knowledge` or `sound`. Of these edits, 98.5% were located and applied by exact match, and every edit was marked and kept reversible. Every conversation containing words was corrected except six. The pass used 57.3 million input and 19.4 million output tokens. It cost US$154.13 at the batch price (US$0.07 per audio hour), or US$308.26 at the standard price, at which Google billed the batches completed by October 21.

2.4 Speakers

Each speech segment was represented by a voiceprint, an embedding from the frozen (not retrained) WeSpeaker ResNet-34 model (Wang et al., 2023), normalized by room type. A logistic-regression classifier assigned each voiceprint to one of fourteen voices (Nixon, Kissinger, Haldeman, Ehrlichman, Colson, Dean,

Ziegler, Mitchell, Rogers, Shultz, Haig, Woods, Connally, Butterfield) or to "other". It was trained on 24,656 clips from 902 conversations: WSPF speaker turns, two-person recordings on which Nixon was confidently ruled out (bringing in Kissinger, Haig, Rogers and others), and the author's 22 marks. Each Whisper segment of at least 1.5 seconds was scored only against the people NARA listed as present and was named when the classifier gave a person a probability of 0.7 or more.

2.5 Evaluation

Each machine transcript and its reference were normalized with Whisper's English normalizer (Radford et al., 2023), cut to the stretch they shared and aligned word by word. Word recall, $1 - (S + D)/N$, measured the share of the N reference words that were neither substituted (S) nor deleted (D); the word error rate (WER), $(S + D + I)/N$, also counted inserted words (I). Name recall and name precision were computed over the names in the reference, taken to be words capitalized mid-sentence at least twice and in at least 80% of their uses; a wrong name absent from the reference was not counted. Recall was also measured on Nixon's turns and on replies of four words or fewer. On the Test set, 95% intervals were obtained by bootstrap resampling of conversations, paired between layers. The Test set informed no decision and shared no conversation with the Train set. Of 231 FRUS conversations, 66 were excluded: 44 were not in the corpus (43 released only within whole-tape files, one released after the conversation list was built), 9 had been heard, labeled or corrected by the author, 7 were parallel recordings of excluded conversations, and 6 had audio holding a different conversation. Because FRUS text could have been in the correction model's training data, the merged layer, which no language model touched, was scored alongside the corrected layer.

3 Results

3.1 Accuracy

Fine-tuning worked. On the six train room recordings, none of which were in the fine-tuning data, fine-tuning raised Parakeet's word recall from 63.8% to 72.3%, a paired gain of 8.5 points (95% interval 5.0 to 10.8; Figure 2c). Recall on Nixon's own words rose from 49.3% to 64.9% and on short replies from 32.2% to 46.7%, while WER fell from 39.0% to 31.3%.

The corpus beat Whisper on held-out data. On the 165 FRUS conversations, the corrected layer recovered 76.1% of reference words, against 69.2% for Whisper alone (Table 1, Figure 2). Of the 6.9-point gain (interval 6.3 to 7.7), 4.5 points came from fine-tuned Parakeet and 2.4 from correction. The corpus was more accurate than Whisper on 160 of the 165 conversations. The gap between the two recognizers was largest on Nixon's own words, which Parakeet v5 recovered at 71.6% and Whisper at 63.6%, and on short replies, at 50.8% and 36.0%. By location, the corrected layer recovered 85.6% of words at Camp David, 83.7% on the telephone, 76.6% in the Oval Office, 76.5% in the Cabinet Room and 49.3% in the EOB (Figure 2d).

Correction supplied the names. Parakeet v5 recovered 60.2% of names and Whisper 60.0%. Correction raised name recall by 19.8 points, to 80.1%, with no loss of precision: of the reference names the corrected layer displayed, 90.3% were right, against 89.0% for Whisper alone.

Table 1: Accuracy of each corpus layer (%). Train: the five train room recordings in the corpus (48,189 reference words). Test: 165 held-out FRUS conversations (351,883 reference words), with 95% bootstrap intervals over conversations in brackets.

	Train (five room recordings)			Test (165 FRUS conversations)		
Layer	Word recall	WER	Names	Word recall	WER	Names
Whisper large-v3 alone	57.2	45.8	48.8	69.2 [67.4, 70.7]	42.0	60.0 [57.0, 62.8]
Parakeet v5 alone	72.4	31.3	73.0	73.7 [72.0, 75.1]	38.1	60.2 [57.2, 63.2]
Merged (no language model)	72.5	31.1	73.1	73.7 [72.1, 75.2]	37.9	60.3 [57.2, 63.2]
Corrected	**73.9**	**29.6**	**81.2**	**76.1** [74.5, 77.6]	**35.1**	**80.1** [77.3, 82.5]

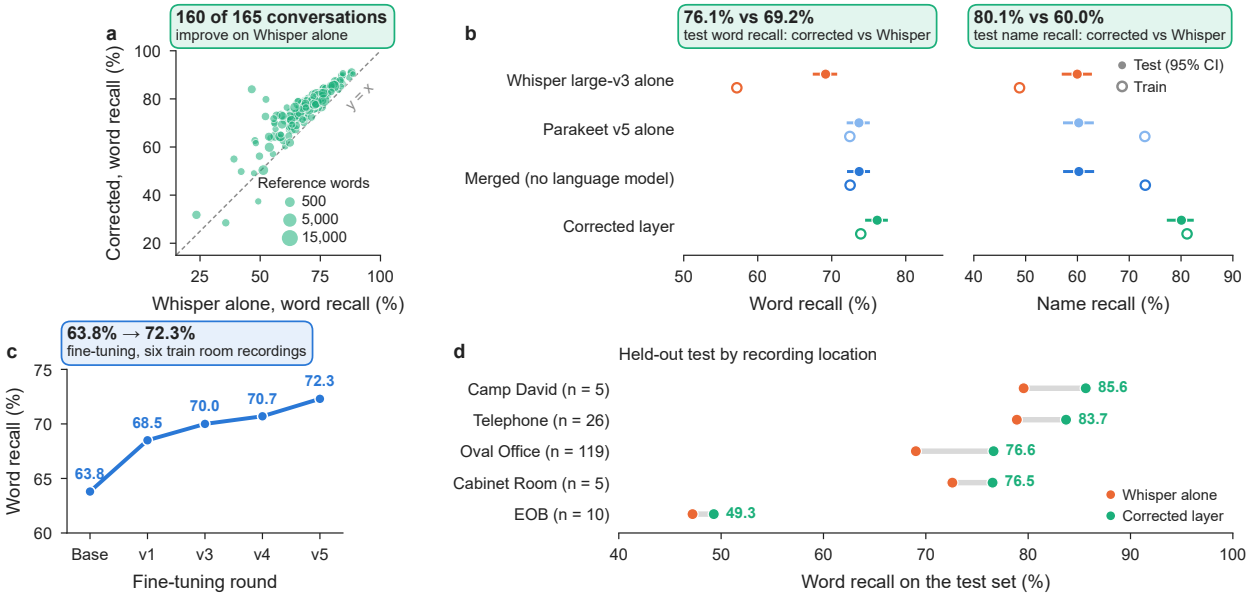


Figure 2: Accuracy. **a**, Word recall of each of the 165 held-out FRUS conversations, Whisper alone against the corrected layer (marker size: reference words). **b**, Word and name recall of each layer on the Test set (filled, 95% intervals) and the five train room recordings in the corpus (open). **c**, Fine-tuning: pooled word recall on the six train room recordings, by round. **d**, Test word recall by recording location, Whisper alone and corrected.

3.2 Artifacts of single-model transcription

Whisper alone looped, dropped speech and wrote over silence. Its output (Figure 3) contained 2,613 repetition loops, the same words written over and over, in 1,463 conversations (55,140 words); Parakeet v5 produced 105 loops and the released corrected layer 265. When run in batches with WhisperX (Bain et al., 2023), the configuration of earlier machine transcriptions of the tapes (Fine, 2026), Whisper also discarded speech: wherever WhisperX wrote nothing for 20 seconds or more, Parakeet still transcribed speech. These 3,973 stretches held 5.2% of Parakeet's words in those conversations, and in 2,869 of them Parakeet wrote at least 60 words a minute, a rate that only real speech produces. Over NARA's muted withdrawals, Whisper wrote 1,882 words in 422 conversations, typically "you" every 30 seconds, while Parakeet wrote 131 words over the same silence. Across the corpus, Whisper also wrote 987 stock phrases such as "thank you for watching". No word inside a muted passage was kept in the release. Each recognizer also heard words that the other missed, and both were therefore kept in the corpus.

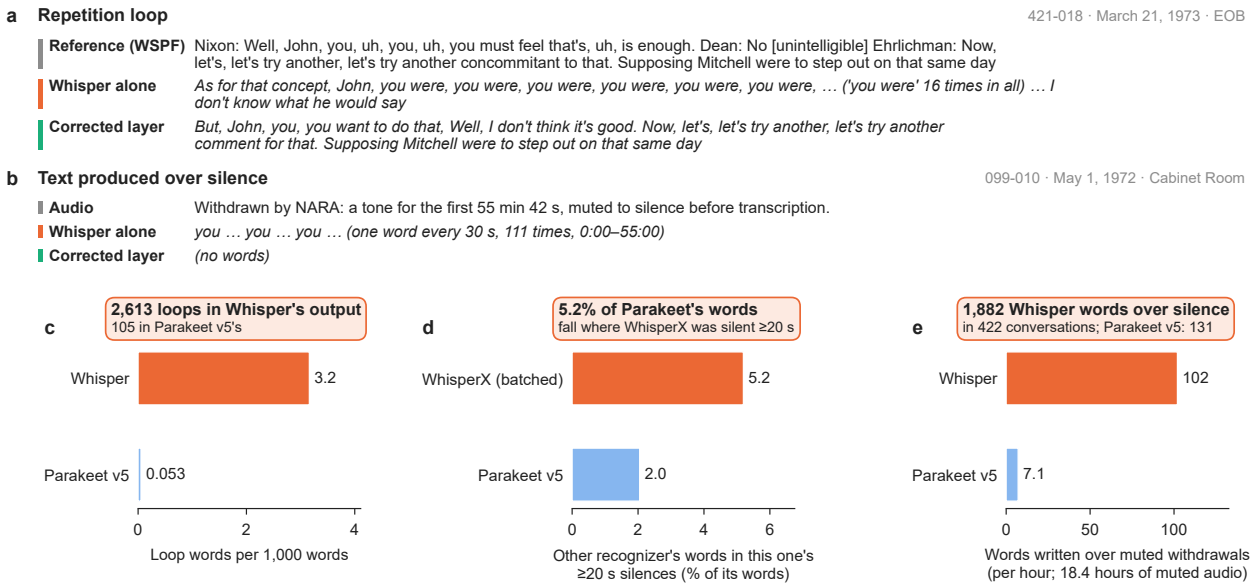


Figure 3: Artifacts of Whisper-only transcription, from the corpus's own Whisper layer. **a**, A repetition loop (421-018): Whisper repeated "you were" sixteen times and lost Ehrlichman's turn; the corrected layer recovered it. **b**, Text over silence (099-010): across a 55-minute withdrawal, muted before transcription, Whisper wrote "you" every 30 seconds; the corrected layer held no words. **c**, Loop words per 1,000 words, whole corpus. **d**, Dropped speech: the other recognizer's words that fell in this one's silences of 20 seconds or more (WhisperX: the batched fifth). **e**, Words written over muted withdrawals per hour of muted audio.

3.3 What the second reading and the correction added

On the Test set, at least one recognizer was right on 79.5% of reference words, 6.0 points more than Parakeet alone on the same words, and correction recovered 2.4 of those points.

The second reading drove the gain. On the five train room recordings in the corpus, correction gained 405 reference words when given Parakeet's text alone, 564 when Whisper's readings were added and 712 when low-confidence words were also marked; before the corpus run, 340 edits helped and 20 hurt.

Same-day context corrected across recordings. On April 9, 1971, Parakeet's *"want Dr. Burns, uh, Marjorie Burns? I'm pleased to call"* (476-021) was corrected to *"Find Dr. Burns, Arthur Burns, I placed a call"*, the words of the telephone recording of the same moment (001-061); the author confirmed both by ear.

Edits cited their evidence. Of the corrected names in the Test set, 42.8% cited Whisper's reading and 34.8% the subject log; only 4.6% cited the model's own knowledge.

3.4 Speakers

On the 130 Test conversations outside the classifier's training data (220,664 aligned words), the classifier named the speaker of 68.4% of words, and 92.4% of named words carried the speaker named in the FRUS transcript. For the fourteen voices it was trained on, the figure was 95.9%: 94.9% for Nixon over 78,305 named words and 97.5% for Kissinger over 50,743 (Figure 4c). Its cross-validated balanced accuracy, the average of its accuracy on each voice, was 87.2%.

Table 2: The Nixon Corpus, release 1.1.

Conversations	19,426 (February 16, 1971 – July 12, 1973), 17,343 with words
Audio	2,206 hours: 6,548 telephone, 6,224 Oval Office, 4,435 EOB, 2,003 Camp David, 216 Cabinet Room
Words	18,994,348, each with start and end time, source and confidence; every model edit marked and reversible
Speakers	63% of words named; of 1,461 hours of scored speech, 1,023 named (Nixon 637, Haldeman 99, Kissinger 79)
Context	NARA date and location for every conversation, participants for all but 35, subject log for 94%
Format	JSON (gzip), plain text, metadata table, raw layers, checksums; 776 MB

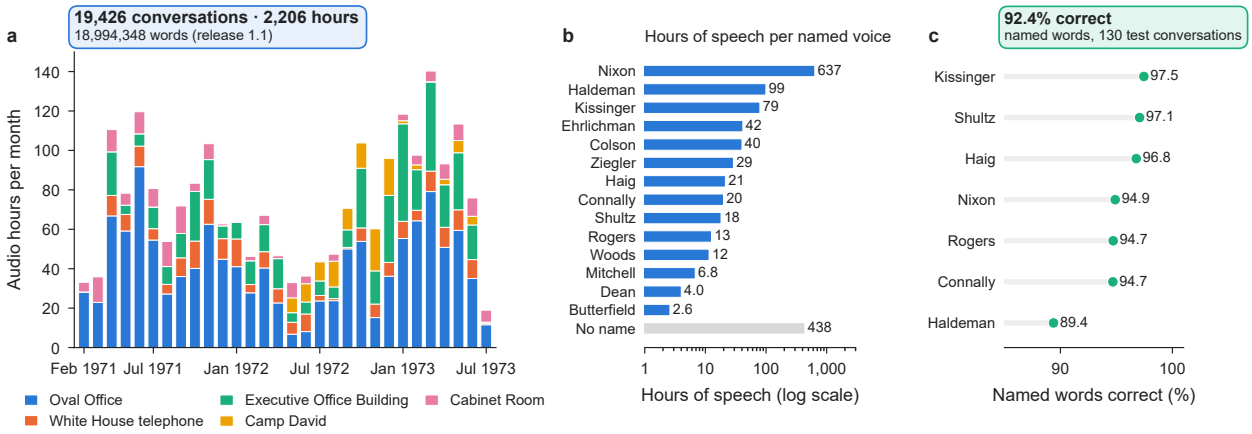


Figure 4: The corpus. **a**, Audio hours per month by recording location, February 1971 to July 1973. **b**, Hours of speech per named voice (segments named at probability 0.7 or more; log scale), and the 438 scored hours without a name. **c**, Named words correct on the 130 held-out Test conversations outside the classifier's training data, for voices with at least 200 named words.

4 The corpus

Release 1.1 of the corpus (Table 2, Figure 4) was packaged as one machine-readable JSON record per conversation. Each record contained NARA's metadata and subject log, the audio sources, the withdrawn spans, the speaker segments, and every word with its timing, speaker, confidence and source. The source recorded whether a word came from Parakeet, both recognizers, a Whisper fill, a model edit or a correction by ear. Each edit was stored with the words it replaced and the reason for the change. Plain-text transcripts and the raw output of each recognizer were also included. A browsing application, included in the repository, was built to present the corpus by day, conversation and person and to play each transcript in sync with the audio; no language model wrote its text.

5 Three passages

The following quotations were taken verbatim from the released corrected transcript, with its punctuation and standard capitalization and spelling, and were checked by the author against the audio.

5.1 The corpus matched a participant's transcript: 871-004, March 7, 1973

With L. Patrick Gray's Senate confirmation as FBI director going badly (Congressional Research Service, 2018), Nixon told White House Counsel John Dean that he wanted friendly senators to question Gray under oath about what J. Edgar Hoover had told him. Dean later printed the passage from his own transcription (Dean, 2014, Conversation No. 871-4):

> **Dean (2014):** "I'm going to make him lie, because I think Gray's not handling himself well. What he said yesterday with regard to the fact that this 'jolly well' bullshit … that it was highly improper for you to be present …"
>
> **This corpus (5:26–5:48):** *"I told him that Hoover had told me, and I'm going to make him lie because I think Gray's not handling himself well. What he said yesterday with regard to the fact that, it was, that, you know, this jolly well bullshit and all that sort of thing that you were, that it was highly improper for you to be present."*

These were the recognizers' own words; the correction step made no edit here. Over the surrounding minutes, the merged layer shared 95 of its 1,288 five-word sequences with Dean's text, against 10 to 20 for unrelated stretches. Dean printed only what he selected, whereas the corpus holds the whole conversation, including Nixon's question at 4:55 about the FBI's 1968 *"electronic surveillance of Vice President Agnew's phone when he was in Phoenix"*. The documented check concerned calls from Albuquerque, not Phoenix (Germany et al., 2014; Hughes, 2014): the recognizer heard Nixon correctly, and the slip was his.

5.2 The cover story in Nixon's words: 916-019, May 11, 1973

Eleven days after Haldeman and John Ehrlichman resigned, and days after the CIA's deputy director, Vernon Walters, testified that the two had sought the CIA's help in stopping an FBI investigation (Brinkley and Nichter, 2015), Nixon met Haldeman and Alexander Haig, Haldeman's successor as chief of staff, in the Oval Office. At 10:20 Nixon, whose voice the author confirmed by listening, told Haldeman how Ehrlichman's account of the June 23, 1972 meeting with the CIA's leaders should go:

> *"Now, my point is, John's story should be perfect, should be, and I hope you talked to John. John's story is perfect. John should say his only concern was the whole CIA involvement and so forth and wanted to be sure that the, this investigation of Watergate didn't get out, as he put it, didn't get out of hand. […] Getting out of hand had to do with exactly what's now happening. Yeah. Which is unloading all this Pentagon papers. Investigation type stuff, the leak, the plumbers operation type stuff, the, Exactly. Wiretapping of the reporters and, and staff members type stuff."*

"John" was Ehrlichman: the conversation placed *"Ehrlichman and me"* at the meeting and said that *"Dean was not there"* and that *"Ehrlichman did do all the talking"*. Parakeet transcribed these words where Whisper wrote nothing, and the correction step did not touch them. At 13:00 Nixon named his concern, and Haldeman answered (the author corrected *"faint"* to *"paint a"* by ear):

> **Nixon:** *"I, I don't need to paint a picture. I just want to be sure that John doesn't indicate that the President directed him to get Walters to get Gray to turn this off and set up a phony story"*
>
> **Haldeman:** *"John would never indicate that. […] What the President tells us to do is not, not public, public consumption."*

Table 3: Further leads, with prior-publication checks. Three other candidates were already in print (Dean, 2001, 2014).

Conversation	Date, place	Passage	Prior publication
789-006	Sept. 30, 1972, Oval Office	*"super apathy among the blacks"*, called *"just exactly where you want it"* before the election	Not in the registry or Brinkley and Nichter (2014)
233-010	Dec. 6, 1972, Camp David	*"make Sweden the black post. I want a black for Sweden"*	Not in Brinkley and Nichter (2014); in earlier machine transcripts (Fine, 2026)
690-011	Mar. 21, 1972, Oval Office	An ambassadorship and a request to *"get at least 150 out of it"*	Not in Brinkley and Nichter (2014); in earlier machine transcripts (Fine, 2026)

It was Nixon who, on June 23, 1972, told Haldeman that the CIA's leaders "should call the FBI in and say that we wish for the country, don't go any further into this case, period!" (Watergate Special Prosecution Force, 1974). No published transcript of this exchange[1] was found in the registry, in either volume by Brinkley and Nichter, in Dean (2014) or on the Nixon Library's and Miller Center's pages for the conversation. Kutler (1997) and the Presidential Recordings Digital Edition were not checked.

5.3 The ordinary record, and further leads

Most of the tapes are routine, and the corpus captures that too. On March 14, 1972, after former Attorney General John Mitchell's testimony on the ITT affair, Nixon telephoned his special counsel Charles Colson and tried out a line, *"this is a smear a day Congress rather than a law day of Congress"*; Colson laughed, and they agreed to use it (021-068). No subject log records such exchanges. Further leads were surfaced by ranking passages from conversations outside the registry by their politically charged terms; the author listened to each lead and confirmed it (Table 3).

6 Limitations and conclusion

Limitations. The corpus is a machine transcript and a finding aid. It recovered 76.1% of test reference words, and 49.3% in the EOB, so every quotation should be checked against the audio, to which every word is linked. FRUS editors chose the test excerpts; a test set drawn at random from the corpus would measure the whole record. Correction depends on a proprietary model, but its edits are marked and reversible. Voices without a class of their own were sometimes given a trained speaker's name (16.9% of their test words). Withdrawn passages stay muted, and passages about private persons call for care in republication.

Conclusion. Two recognizers that err differently, reconciled against the archive's own records, turn the released Nixon tapes into a searchable, speaker-labeled text of 19.0 million words. On held-out conversations this text recovered 76.1% of words and 80.1% of names, against 69.2% and 60.0% for Whisper alone. It tells a historian where to listen, with measured accuracy, word timings and speaker labels that neither the subject logs nor earlier machine transcripts provide.

[1] Easy Nixon's unreviewed machine transcript of 916-019 (Fine, 2026) did not contain it.

Declarations

Data, code and model availability. The corpus (release 1.1) is archived at `https://doi.org/10.5281/zenodo.23657174`, with a data card that documents its format, coverage and limitations. The code for the pipeline, the evaluation and the browsing tool, the scripts that rebuild the public-domain reference transcripts, and the correction instructions (`src/argon.py`) are at `https://github.com/akoopie/nixon-corpus`. The fine-tuned recognizer is at `https://huggingface.co/akoopie/nixon-parakeet-tdt-0.6b-v5`. The corpus viewer is available at `https://nixoncorpus.org`, where the transcripts can be read and searched with the audio playing alongside.

License and citation. The transcripts, annotations and model weights are released under CC BY 4.0; attribution should take the form of a citation of this paper and of the release. The model derives from NVIDIA's Parakeet-TDT-0.6B-v3 (CC BY 4.0; NVIDIA, 2025). The code is MIT-licensed. The recordings, NARA's metadata and the government transcripts are public domain; please credit NARA and the Richard Nixon Presidential Library and Museum as their source.

Use of AI tools. AI systems were used as components of the pipeline (the speech recognizers, the speaker-embedding network and the correction model, Gemini 4 Argon, as reported above) and as assistants: the software was written, the analyses were run and the manuscript was drafted with the help of an AI coding assistant (Claude, Anthropic). The author designed the study, directed all analyses, made the research decisions, performed all annotation and listening, and takes full responsibility for the content of this paper.

Acknowledgments. The author thanks the National Archives and Records Administration and the Richard Nixon Presidential Library and Museum, whose archivists preserved, reviewed, described and released the recordings, and whose subject logs and conversation manifest made this work possible.

Competing interests and funding. The author declares no competing interests. The work received no external funding; computing costs were paid by the author.

References

Alexei Baevski, Henry Zhou, Abdelrahman Mohamed, and Michael Auli. wav2vec 2.0: A framework for self-supervised learning of speech representations. In *Advances in Neural Information Processing Systems 33*, 2020. URL `https://arxiv.org/abs/2006.11477`.

Max Bain, Jaesung Huh, Tengda Han, and Andrew Zisserman. WhisperX: Time-accurate speech transcription of long-form audio. In *Proceedings of Interspeech 2023*, pages 4489–4493, 2023. doi: 10.21437/Interspeech.2023-78.

Mateusz Barański, Jan Jasiński, Julitta Bartolewska, Stanisław Kacprzak, Marcin Witkowski, and Konrad Kowalczyk. Investigation of Whisper ASR hallucinations induced by non-speech audio. In *ICASSP 2025 – IEEE International Conference on Acoustics, Speech and Signal Processing*, 2025. URL `https://arxiv.org/abs/2501.11378`.

Hervé Bredin. pyannote.audio 2.1 speaker diarization pipeline: Principle, benchmark, and recipe. In *Proceedings of Interspeech 2023*, pages 1983–1987, 2023. doi: 10.21437/Interspeech.2023-105.

Douglas Brinkley and Luke A. Nichter. *The Nixon Tapes: 1971–1972*. Houghton Mifflin Harcourt, Boston, 2014.

Douglas Brinkley and Luke A. Nichter. *The Nixon Tapes: 1973*. Houghton Mifflin Harcourt, Boston, 2015.

Chen Chen, Yuchen Hu, Chao-Han Huck Yang, Sabato Marco Siniscalchi, Pin-Yu Chen, and Eng Siong Chng. HyPoradise: An open baseline for generative speech recognition with large language models. In *Advances in Neural Information Processing Systems 36, Datasets and Benchmarks Track*, 2023. URL https://arxiv.org/abs/2309.15701.

Congressional Research Service. FBI director nominations, 1973–2017. Report R44842. https://www.congress.gov/crs-product/R44842, 2018.

John W. Dean. *The Rehnquist Choice: The Untold Story of the Nixon Appointment That Redefined the Supreme Court*. Free Press, New York, 2001.

John W. Dean. *The Nixon Defense: What He Knew and When He Knew It*. Viking, New York, 2014.

Niall Ferguson and Luke Nichter. The Nixon tapes, unlocked. *The Free Press*, September 30. https://www.thefp.com/p/the-nixon-tapes-unlocked-niall-ferguson-luke-nichter, 2026.

Todd Fine. Easy Nixon: Accessible archive of the Richard M. Nixon white house tapes. https://www.easynixon.org/, 2026. Accessed October 21, 2026.

Jonathan G. Fiscus. A post-processing system to yield reduced word error rates: Recognizer output voting error reduction (ROVER). In *1997 IEEE Workshop on Automatic Speech Recognition and Understanding Proceedings*, pages 347–354, 1997. doi: 10.1109/ASRU.1997.659110.

Kent B. Germany, Nicole Hemmer, and Ken Hughes. Lyndon B. Johnson and Cartha D. "Deke" DeLoach on 13 november 1968, conversation WH6811-04-13733. In *Presidential Recordings Digital Edition: Johnson Telephone Tapes, 1968*. University of Virginia Press, Charlottesville, 2014. URL https://prde.upress.virginia.edu/conversations/4006140.

Google DeepMind. Gemini 4 argon: Model card. https://deepmind.google/models/model-cards/gemini-4-argon/, 2026. Published October 13, 2026.

H. R. Haldeman. *The Haldeman Diaries: Inside the Nixon White House*. G. P. Putnam's Sons, New York, 1994.

Hoover Institution. Hoover applied history working group announces landmark transcription of Nixon white house tapes. Press release, September 30. https://www.hoover.org/news/hoover-applied-history-working-group-announces-landmark-transcription-nixon-white-house-tapes, 2026.

Ken Hughes. *Chasing Shadows: The Nixon Tapes, the Chennault Affair, and the Origins of Watergate*. University of Virginia Press, Charlottesville, 2014.

Allison Koenecke, Anna Seo Gyeong Choi, Katelyn X. Mei, Hilke Schellmann, and Mona Sloane. Careless Whisper: Speech-to-text hallucination harms. In *Proceedings of the 2024 ACM Conference on Fairness, Accountability, and Transparency*, pages 1672–1681, 2024. doi: 10.1145/3630106.3658996.

Oleksii Kuchaiev, Jason Li, Huyen Nguyen, Oleksii Hrinchuk, Ryan Leary, Boris Ginsburg, Samuel Kriman, Stanislav Beliaev, Vitaly Lavrukhin, Jack Cook, Patrice Castonguay, Mariya Popova, Jocelyn Huang, and Jonathan M. Cohen. NeMo: A toolkit for building AI applications using neural modules. arXiv:1909.09577, 2019. URL https://arxiv.org/abs/1909.09577.

Stanley I. Kutler, editor. *Abuse of Power: The New Nixon Tapes*. Free Press, New York, 1997.

NARA. White house tapes of the Nixon administration, 1971 to 1973. Conversation-level dataset, Data.gov. https://catalog.data.gov/dataset/white-house-tapes-of-the-nixon-administration-1971-to-1973, n.d. CC0; 22,723 conversations.

Nixon Library. H. R. Haldeman diaries collection finding aid (RN-HRHD; NARA identifier 7787364). https://www.nixonlibrary.gov/sites/default/files/2019-03/HRHaldemanDiariesCollectionFindingAid03212018.pdf, 2018.

Nixon Library. White house tapes: Sound recordings of meetings and telephone conversations of the Nixon administration, 2/16/1971–7/18/1973 (RN-WHT; NARA ID 597542), complete finding aid. https://nixonlibrary.gov/sites/default/files/2019-08/White%20House%20Tapes%20Complete%20Finding%20Aid.pdf, 2019.

Nixon Library. White house tapes: Tape subject logs. https://www.nixonlibrary.gov/white-house-tapes, n.d.

NVIDIA. parakeet-tdt-0.6b-v3. Model card, Hugging Face. https://huggingface.co/nvidia/parakeet-tdt-0.6b-v3, 2025. Released August 14, 2025; CC BY 4.0.

Presidential Recordings Digital Edition. The presidential recordings of Richard M. Nixon. Presidential Recordings Digital Edition. Charlottesville: University of Virginia Press. https://prde.upress.virginia.edu/content/nixon, n.d.

Alec Radford, Jong Wook Kim, Tao Xu, Greg Brockman, Christine McLeavey, and Ilya Sutskever. Robust speech recognition via large-scale weak supervision. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 28492–28518, 2023. URL https://proceedings.mlr.press/v202/radford23a.html.

U.S. Department of State. Foreign relations of the United States, 1969–1976. Washington, DC: Government Publishing Office. Digital edition, https://history.state.gov/historicaldocuments; TEI source, https://github.com/HistoryAtState/frus, n.d. Office of the Historian.

U.S. House of Representatives, Committee on the Judiciary. *Transcripts of Eight Recorded Presidential Conversations*. U.S. Government Printing Office, Washington, DC, 1974. 93rd Congress, 2nd session.

Hongji Wang, Chengdong Liang, Shuai Wang, Zhengyang Chen, Binbin Zhang, Xu Xiang, Yanlei Deng, and Yanmin Qian. Wespeaker: A research and production oriented speaker embedding learning toolkit. In *ICASSP 2023 – IEEE International Conference on Acoustics, Speech and Signal Processing*, pages 1–5, 2023. doi: 10.1109/ICASSP49357.2023.10096626.

Watergate Special Prosecution Force. Transcripts of Nixon white house tape recordings. Records of the Watergate Special Prosecution Force, Record Group 460, National Archives and Records Administration; per-conversation copies at https://www.nixonlibrary.gov/white-house-tapes, 1973–1975.

Watergate Special Prosecution Force. Transcript of the conversation of June 23, 1972 between President Nixon and H. R. Haldeman (conversation 741-002), trial exhibit 1, *United States v. Mitchell et al.* https://www.nixonlibrary.gov/sites/default/files/forresearchers/find/tapes/watergate/trial/exhibit_01.pdf, 1974.

Hainan Xu, Fei Jia, Somshubra Majumdar, He Huang, Shinji Watanabe, and Boris Ginsburg. Efficient sequence transduction by jointly predicting tokens and durations. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, 2023. URL https://proceedings.mlr.press/v202/xu23g.html.